FRUCTIFY

Cultivate growth, engineer success.

Katana AI for Clients Primer

VERSION 3 · CLIENT-FACING OFFERING PLAYBOOK

Prepared for: Drew Fitchett, Founder — Katana

By: Fructify LLC

Date: April 2026

Cultivate growth, engineer success.

© 2026 Fructify LLC · Prepared in partnership with Drew Fitchett / Katana

Table of Contents

01 What We're Solving For

02 Questions Clients Are Already Asking

03 Who's Really Asking (Role by Role)

04 Vertical Patterns: Where AI Is Hitting First

05 Katana's AI-for-Clients Service Offerings

06 Integrating AI into Client Systems

07 Integration Approach: MCP vs API vs CLI vs Platform Plug-ins

08 Productized AI Building Blocks

09 Vendor Landscape & Licensing

10 Enterprise-Grade & Verticalized AI

11 Data Requirements & Readiness

12 Security, Risk & Governance (Including AI Policy)

13 The 4D Framework: Scoring & Delivering AI Capability

14 Environment Setup & Standardization

15 Pricing, Packaging & Commercial Model (Open Questions)

16 Technical Build Specs & the 30-Day Action Plan

01 What We're Solving For

Katana's SMB and mid-market prospects are asking **AI questions their current MSP can't answer**. They are fielding AI pitches from every direction — Microsoft, Google, industry SaaS vendors, boutique consultants — and they want someone they already trust to cut through the noise.

This primer is the offer-side playbook: **what Katana sells, to whom, how it fits their stack, and what guardrails wrap it**. It is the client-facing pillar of Katana's AI enablement practice. The internal engineering playbook (how Katana's own team uses AI to run the MSP) is a separate document.

> **POSITIONING**
>
> Katana is the **AI-enabled MSP**. Competitors bolted AI on. Katana designed the service with AI, guardrails, and governance baked in from day one.

Why now

- **Client demand has flipped.** Two years ago prospects were skeptical. Today they are actively requesting AI rollouts, Copilot enablement, and governance help.
- **Regulatory floor is rising.** NIST AI RMF, ISO/IEC 42001, state-level AI laws (Colorado SB 205), and sector-specific rules (HIPAA, FINRA, PCI) all touch AI usage.
- **The MSP commoditization trap is real.** Break/fix and managed endpoints are a race to the bottom. AI enablement is a defensible, advisory layer above the commodity.

02 Questions Clients Are Already Asking

When Drew or an account manager walks into a QBR, these are the questions that actually come up — organized by the business outcome the client is really after.

Client Question	What They Really Mean	Katana's One-Line Response
"Should we be using ChatGPT?"	Are we behind? Are our competitors using this?	Probably yes, but on a licensed business tier with your data protected — we'll show you which tier fits and what your team can do with it.
"Is Copilot worth it?"	My Microsoft rep is pushing it. Should I buy?	It depends on your M365 tier, your data hygiene, and how your team actually works — we'll run a 4D readiness check before you commit the licenses.
"Can AI read our email/ contracts/tickets?"	I want AI that knows my business without me retyping it.	Yes — through secure retrieval over your existing data sources (SharePoint, email, ticketing, file shares) with access controls that mirror your permissions.
"Will AI replace my people?"	I'm worried. Should I be laying off or hiring differently?	AI amplifies the staff you have — we help you delegate the right tasks and keep humans in the loop on the rest.
"What's our AI policy?"	Our board/insurer/auditor asked and I don't have one.	We'll draft one based on NIST AI RMF and your industry, then roll it out to staff with training.
"Can you make our SaaS 'smart'?"	My CRM/ERP/PSA is clunky. Can AI help?	Often yes — either through the vendor's native AI (Salesforce Einstein, HubSpot Breeze, NetSuite AI) or a custom agent that integrates with it.
"How do I stop shadow AI?"	I know my team is pasting data into ChatGPT. That terrifies me.	Give them a sanctioned tool, an AUP, and monitoring — banning it is harder and less effective.
"Is my data safe?"	What are these models doing with my inputs?	On business-tier products (Copilot, ChatGPT Enterprise, Gemini for Workspace), your prompts and outputs are not used to train the public model — we'll document this per tool.

03 Who's Really Asking (Role by Role)

Different roles want different things from AI. A good AI conversation adjusts to the seat the client is sitting in.

Role	Their Pain	AI Value They Will Buy
CEO / Owner	Is my business falling behind?	Strategic AI roadmap; vertical-specific wins; productivity lift narrative they can take to the board or partners.
COO / Ops Lead	My people are drowning in repetitive work.	Workflow automation, intake triage, meeting summarization, document generation.
CFO / Controller	License waste; ROI pressure on any new spend.	Utilization reporting on existing AI seats; ROI model before buying more; risk-adjusted business case.
IT Director	AI is landing in every tool — I can't keep up.	Rationalized stack; AI platform strategy; configuration standards; a partner to offload the evaluation work.
Security Officer	Shadow AI; data leakage; DLP gaps.	AUP; DLP posture review; sanctioned-tool rollout; audit trail; policy mapped to NIST AI RMF.
Head of Sales / Marketing	Content velocity; CRM data decay.	Copilot + CRM-native AI (Einstein, Breeze); prospect research agents; email/proposal drafting.
HR Director	Policy writing; recruiting; onboarding docs.	Handbook updates; policy drafting; resume/JD matching; training content generation.
Line-of-business lead	A specific painful process they own.	Vertical or process-specific agent: estimating, intake, compliance check, scheduling, etc.

04 Vertical Patterns: Where AI Is Hitting First

A handful of verticals are doing the heavy lifting in AI adoption right now. These are the best first-call targets for Katana.

Vertical	High-Value AI Use Cases	Regulatory / Data Notes	Why Katana Wins Here
Professional services (law, accounting, consulting)	Contract review, document drafting, time capture, research, client Q&A over knowledge base.	Client confidentiality; bar/AICPA rules; retention and privilege.	Clio/MyCase/QBO integration; data minimization; on-shore LLM options where required.
Healthcare & dental (SMB practices)	Charting assist, scheduling, intake summarization, patient messaging, coding.	HIPAA; BAA required on every AI tool touching PHI.	MSP security posture + BAAs already in motion; HIPAA-aware prompt/AUP.
Construction & trades	Estimating, takeoff assistance, RFI responses, safety doc generation, scheduling.	Less regulated; but IP on bid content and client relationships matters.	Procore/STACK/Buildertrend integration; Copilot for field staff.
Financial services (RIA, wealth, CPA)	Client meeting prep, portfolio narrative, compliance review, email triage.	FINRA, SEC, state insurance regs; supervision rules on communications.	Supervised AI with logging; compliance-ready AUP; Redtail/ Wealthbox/QBO context.
Manufacturing & distribution	Predictive maintenance, SOP generation, RFP response, ERP conversational interfaces.	IP protection; shop-floor data sometimes unstructured.	NetSuite/Dynamics connectors; ability to sit next to OT safely.
Nonprofits & associations	Grant writing, donor communications, board summaries, event planning.	Donor data privacy; 501(c)(3) communication rules.	Low-cost AI bundles; grant-ready templates; M365/Google integration.

05 Katana’s AI-for-Clients Service Offerings

PRICING & PACKAGING NOTE

The “Cadence” column below describes the **shape of delivery** (one-time assessment, recurring, embedded, etc.) — **not** a price or commercial model. Dollar amounts, retainer structures, and margin decisions are Drew’s call and will be set separately in the commercial model (Section 15).

Offering	What's in It	Who Buys It	Cadence
AI Readiness Assessment	4D-framework scorecard, current-state stack review, shadow-AI scan, top 5 use-case recommendations, draft AUP.	Any new prospect; any client without an AI policy.	One-time engagement with a defined deliverable.
Copilot / Workspace AI Enablement	License mapping, prerequisite data hygiene, security baseline, pilot cohort, training, adoption reporting.	M365 or Google Workspace clients evaluating or already licensed for Copilot / Gemini.	Project engagement, then ongoing adoption support.
AI Policy & Governance	AUP drafting, NIST AI RMF mapping, tool inventory, DLP posture review, training rollout.	Any regulated or risk-averse client; any client with a board/ insurer/auditor asking.	One-time build plus ongoing review cycle.
Custom AI Agents / Integrations	Scoped agent builds over client data (intake triage, estimating, Q&A over knowledge base, sales research).	Clients with a specific painful workflow and a willingness to invest beyond licensing.	Scoped project with option for ongoing maintenance / iteration.
vCAIO (virtual Chief AI Officer)	Quarterly roadmap, license optimization, governance oversight, new use-case intake.	Mid-market clients and professional firms with ongoing AI scope.	Ongoing advisory cadence alongside Katana's vCIO service.
AI-Ready Managed Services	Standard Katana managed stack + AI-enabled helpdesk + baseline Copilot/Gemini + AUP.	All Katana managed clients — this is the default, not an upsell.	Embedded in the managed services relationship.

How to sell each

- **Readiness Assessment** is the *foot-in-the-door*. It almost always surfaces downstream work. Lead with it when a prospect is AI-curious but unclear.
- **Copilot / Workspace enablement** is the *immediate-hit* offering — clients already have the license or are being pitched it. Katana owns the rollout end-to-end.
- **AI Policy** is the *compliance lever*. Insurance renewals, SOC 2 scope, board asks — these create urgency and give Katana a seat at the table.
- **Custom agents** are the *high-margin, hard-to-commoditize* work. Build a small library of repeatable agent patterns (intake, estimating, Q&A) and sell variants.

- **vCAIO** creates *ongoing advisory revenue* and protects the account from competitors pitching point solutions.
- **AI-Ready Managed Services** is the *brand moat*. Every Katana client gets baseline AI — that's the story.

06 Integrating AI into Client Systems

AI value explodes when it reaches client data. The integration pattern varies by the system. Katana's job is to pick the right pattern, deliver it safely, and maintain it.

System Type	Examples	Integration Shape	Katana's Role
Productivity suite	M365, Google Workspace	Native AI (Copilot, Gemini) + tenant-level config.	Licensing, rollout, training, adoption reporting.
CRM	Salesforce, HubSpot, Zoho, Dynamics	Native AI (Einstein, Breeze, Copilot for Sales) or custom agent over CRM data.	Field/object mapping, data hygiene, workflow design.
ERP / Accounting	NetSuite, Dynamics, QBO, Sage Intacct	Native AI where available; conversational agent over GL/reports otherwise.	Data-model literacy, access control, compliance review.
PSA / Ticketing	ConnectWise, Autotask, HubSpot Service	Native ticket summarization + Katana internal agents; client-facing triage agent where appropriate.	Deflection design, escalation rules, QA.
Document stores	SharePoint, Google Drive, Dropbox, Box	RAG (retrieval-augmented generation) over indexed content; permissions-aware.	Permission mirroring, index hygiene, DLP policy.
Communication	Teams, Slack, Gmail, Outlook	Native copilots + custom bots where a workflow needs it.	Bot deployment, audit, prompt/response logging.
Vertical SaaS	Clio, Procore, Redtail, Epic, Salesforce vertical clouds	Vendor-provided AI feature set; custom integrations where gaps exist.	Vendor liaison, configuration, user enablement.

PRINCIPLE

Start with the **vendor's native AI** before building custom. Custom agents should solve problems vendors can't, not duplicate what a license already delivers.

07 Integration Approach: MCP vs API vs CLI vs Platform Plug-ins

Under the hood, every AI integration is one of four shapes. Knowing which shape to use is a core Katana competency — it affects cost, latency, security posture, and how much of the work a vendor will do for you.

Approach	When to Use	Strengths	Trade-offs
Platform plug-in / native AI	The client's core platform (M365, Salesforce, HubSpot) offers an AI feature that covers the use case.	Fastest to deploy; vendor handles data residency and security.	Limited to vendor's model and scope; pricing set by vendor; feature lag.
MCP (Model Context Protocol)	A Katana-curated agent needs access to multiple client systems and we want a thin standardized connector layer.	Standard protocol; ecosystem of pre-built servers; easy to add/ remove tools.	Adds context surface; loads many tool schemas into context; needs governance on which servers are enabled per client.
Direct API	Production workflow that runs repeatedly at predictable cost and needs tight error handling.	Minimum context cost; most controllable; best for productized agents.	More code; Katana owns the connector maintenance.
CLI / script	One-off data migrations, batch analysis, or engineer-facing internal tools.	Simple, scriptable, auditable.	Not for end-user workflows; no UI.

Katana's default rule

- **Client end-user workflow?** Start with platform-native AI. If the platform covers it, stop.
- **Katana-built agent serving the client?** Use direct API for production paths and MCP only for the connectors where it earns its keep.
- **Internal engineering tool?** CLI or script is fine.

08 Productized AI Building Blocks

Rather than bespoke every time, Katana builds a library of repeatable agent patterns. Each pattern is a pre-solved problem wrapped in branding, security, and a rollout playbook.

Pattern	What It Does	Best Fit	Integrations
Intake triage	Categorizes inbound requests (tickets, emails, leads), routes, drafts initial response.	Professional services, MSPs, healthcare front office.	PSA, CRM, email.
Knowledge Q&A (RAG)	Answers questions over a client's documents with citations.	Any client with a document-heavy knowledge base (SOPs, handbooks, contracts).	SharePoint, Drive, Confluence, Box.
Meeting prep + summary	Pulls context from CRM/email, generates prep docs, summarizes outcomes back.	Sales orgs, professional services.	Calendar, CRM, email, Gong/Granola-style tools.
Proposal / estimate drafter	Generates first-draft proposals or estimates from a structured request.	Professional services, construction, SaaS resellers.	CRM, proposal tool, pricing sheets.
Compliance / policy assistant	Reviews documents against a policy or regulation; flags issues.	Legal, finance, healthcare, HR.	Document stores, policy library.
Sales research agent	Researches a prospect company + contacts; drafts personalized outreach.	B2B sales teams, professional services BD.	CRM, web search, email.
Internal helpdesk bot	Employee self-serve for IT/HR questions; deflects Tier 1 tickets.	Any mid-market client with 100+ employees.	IT doc store, HR handbook, PSA.

09 Vendor Landscape & Licensing

Katana's job is to steer clients through a confusing licensing landscape. The table below shows the current state of the major business-grade AI offerings. **All dollar figures are vendor list prices**.

> **IMPORTANT — THESE ARE VENDOR PRICES, NOT KATANA'S PRICES**
>
> Every price in this section is what the **vendor** charges the client (Microsoft, Anthropic, OpenAI, Google, etc.), shown for reference. **Katana's own fees, margins, and commercial model are set separately by Drew** and are not represented here. See Section 15 for commercial-model open questions.

Vendor / Product	Audience	Vendor List (as of April 2026)	Katana Notes
Microsoft 365 Copilot	M365 E3/E5 customers.	$30 / user / month (annual commit), on top of base M365 license.	Default for any M365 shop; requires data-hygiene prep; best adoption with training.
Copilot for Sales / Service	Sales and service orgs on Dynamics / Salesforce / HubSpot.	$40–$50 / user / month.	Layered on top of M365 Copilot for sales-specific workflows.
ChatGPT Enterprise / Team	Any org wanting a general-purpose AI with enterprise data controls.	Team: $30 / user / month (annual). Enterprise: custom negotiated.	Good fit where M365 Copilot is overkill or the client wants a broad-purpose chat tool; no training on customer data.
Gemini for Google Workspace	Google Workspace customers.	Bundled into Workspace Business/Enterprise (no separate add-on as of Jan 2025).	Default for Google shops; competes directly with Copilot; good for document-heavy workflows.
Anthropic Claude Pro / Team / Enterprise	Org wanting best-in-class reasoning and long context.	Pro: $20 / user / month. Team: $25 / user / month (5 seat minimum). Enterprise: custom.	Katana's preferred general-purpose model; strongest for long-document reasoning and writing-heavy work.
Perplexity Enterprise Pro	Research-heavy roles (BD, consulting, strategy).	$40 / user / month.	Pitch to firms doing a lot of competitive or market research.
Salesforce Einstein / Agentforce	Salesforce customers.	Usage-based (per action / per conversation) + add-on SKUs.	Heavy integration; usually sold/ configured by Salesforce partners — Katana coordinates.
HubSpot Breeze	HubSpot customers.	Included with Pro/Enterprise tiers, with usage credit caps.	Strong fit for SMBs on HubSpot; Katana leads the rollout.
Anthropic API / OpenAI API / Gemini API	Custom Katana-built agents.	Per-token usage pricing; varies by model tier.	Underlying fuel for custom agents; Katana provisions and governs usage.

Vendor pricing changes frequently. Validate the current number at the vendor's site before quoting in any client deliverable.

10 Enterprise-Grade & Verticalized AI

As clients mature, they move from general-purpose AI into enterprise-grade or vertical-specific AI. Katana's job is to know the map.

Category	Representative Offerings	When to Recommend
Enterprise general-purpose	ChatGPT Enterprise, Claude Enterprise, Gemini Enterprise, Copilot with enterprise data boundary.	Mid-market clients who need data-residency, SSO/SCIM, audit logs, admin controls.
Legal	Harvey, Spellbook, CoCounsel (Thomson Reuters), Clio Duo.	Firms doing volume contract review, research, or drafting.
Healthcare	Abridge, Nuance DAX, Suki, Hippocratic AI, Epic's embedded AI.	Clinical practices and hospital systems — with BAA coverage.
Financial services	Hebbia, AlphaSense, Kensho, NVIDIA's banking agents.	Wealth management, IB, research-heavy finance roles.
Sales & marketing	Gong AI, Clay, 11x, Regie.ai, Apollo AI, Jasper.	Sales orgs with clear top-of-funnel or research bottlenecks.
Customer support	Intercom Fin, Zendesk AI, Forethought, Ada.	Clients with a volume support operation and deflection targets.
HR / recruiting	Gem, Paradox, Eightfold, Lattice AI.	Clients with active recruiting or engagement goals.
Engineering / IT	GitHub Copilot, Cursor, Anthropic's Claude Code, Replit Agent.	Clients with internal dev teams.

11 Data Requirements & Readiness

AI is only as good as the data underneath it. A readiness assessment almost always finds the same classes of problems.

Readiness Dimension	What Good Looks Like	Typical Issues Katana Finds
Access & permissions	Least-privilege permissions in SharePoint/Drive/CRM; groups, not individuals; regularly audited.	“Everyone” access on sensitive folders; ex-employees still in groups; AI surfaces content the person shouldn’t see.
Document hygiene	Duplicates cleaned; canonical locations; consistent metadata; archive policy.	5 copies of the same doc at different versions; AI answers from the wrong one.
Structured data quality	CRM and ERP records complete; consistent fields; deduplicated.	Half-empty CRM; AI can’t reliably summarize customer history.
Sensitive-data classification	Labels or DLP rules flag PII/PHI/PCI; AI tools respect those labels.	No classification at all; nothing stops Copilot from surfacing SSNs.
Content lifecycle	Retention rules; stale content archived; policy documents current.	Policies from 2019 still authoritative; AI gives outdated guidance.

SALES ANGLE

Data readiness is **where Katana earns trust before the fun AI part starts**. Every readiness finding is a downstream project opportunity — and a legitimate reason to slow an AI rollout until the posture is right.

12 Security, Risk & Governance (Including AI Policy)

Every AI rollout rides on a security and governance layer. Katana owns this layer — it is the single biggest reason a client should choose an AI-enabled MSP over a point-solution consultant.

Control Area	What Katana Delivers	Framework / Source
Acceptable Use Policy (AUP)	Plain-English policy covering sanctioned tools, prohibited data types, disclosure rules, attribution, review.	NIST AI RMF Govern function; industry templates.
Tool inventory & approval	Living list of sanctioned AI tools, with data-handling notes; process for requesting new tools.	NIST AI RMF Map / Measure.
Data classification & DLP	M365 / Google sensitivity labels; DLP rules that block paste of regulated data into unsanctioned tools.	Vendor DLP tooling; client compliance regime (HIPAA, PCI, etc.).
Access & identity	SSO + MFA on every AI tool; SCIM or just-in-time deprovisioning.	Standard Katana identity baseline.
Logging & audit	Prompt/response logs where regulation requires; admin audit of AI tool usage.	Sector-specific (FINRA supervision, HIPAA audit).
Training & awareness	Short, role-specific AI training; annual refresh; executive briefing.	AUP + industry updates.
Vendor / third-party review	Before a new AI tool is sanctioned: data-handling review, BAA / DPA, model training opt-out, country of processing.	Katana's vendor intake process.
Incident response	Runbook for AI-specific incidents (prompt injection, data leakage via AI, hallucinated output causing harm).	Incorporated into the client's broader IR plan.

13 The 4D Framework: Scoring & Delivering AI Capability

Katana uses Anthropic's AI Fluency 4D framework as the structural spine for every AI engagement. It is the vocabulary Katana uses with clients and the scorecard the readiness assessment produces.

Dimension	Question	What Good Looks Like	Typical Gaps
Delegation	What work should AI do, what should a person do, and how do they hand off?	Clear task map; humans handle judgment; AI handles repetitive drafting / summarization; escalation paths.	Either “AI does everything” fantasy, or nothing delegated at all.
Description	How well does the team describe what they want? Prompts, context, examples, guardrails.	Prompt templates; shared prompt library; role-specific starting points.	Every user reinvents the wheel; inconsistent quality.
Discernment	How well does the team judge AI output? Accuracy, bias, completeness.	QA culture; human-in-the-loop on anything consequential; known-good reference checks.	Copy/paste without review; overconfident outputs making it to clients.
Diligence	What controls sit around AI use? Policy, security, auditability.	AUP; sanctioned tools; DLP; logging; training.	Shadow AI; no AUP; no logging.

How Katana uses it

- **Readiness Assessment** scores the client 1–5 on each D.
- The output is a **4D-scored report, prioritized roadmap, scoped options** for the client to choose from.
- At every QBR, Katana re-scores the client. The trendline is the story.

14 Environment Setup & Standardization

Katana’s default AI-ready client environment. This is the configuration baseline every client inherits when they sign.

Layer	Default	Notes
Identity	Entra ID or Google Workspace with MFA on every account; SSO to every AI tool.	Existing Katana baseline — AI tools inherit it.
Endpoint	Managed endpoints with EDR; browser managed; AI tool installers via MDM or IdP.	No shadow AI tool installs on managed devices.
Data	SharePoint/Drive organized by team with least-privilege permissions; sensitivity labels on regulated content.	Pre-AI data hygiene is a required step before Copilot / agents go live.
AI tool sanctioning	Copilot or Gemini (platform-matched); Claude / ChatGPT Business or Enterprise as secondary where justified.	Documented in the client AUP.
DLP	Rules blocking paste of PII/PHI/PCI into unsanctioned tools; browser-level where needed.	Tuned per industry.
Logging	Admin-level audit on sanctioned AI tools; prompt/response retention where regulation requires.	Fed into Katana's security operations.
Training	Short welcome module on AI use; role-specific modules; annual refresh.	Delivered via existing Katana LMS partner.

15 Pricing, Packaging & Commercial Model (Open Questions)

THIS SECTION IS DELIBERATELY OPEN

Katana's own pricing, margin, and commercial model are Drew's call — not decisions this playbook has made. The items below are the open questions Drew should answer (in partnership with his go-to-market advisors) to lock the commercial model.

Open questions Drew needs to answer

1. **Readiness Assessment** — fixed engagement fee, discounted for existing managed clients, or free loss-leader for qualified prospects?

2. **Copilot / Workspace enablement** — per-seat project fee, cohort-based flat fee, or bundled into a managed-services uplift?
3. **AI Policy & Governance** — one-time engagement, or included in vCAIO / vCIO retainer?
4. **Custom AI agents** — per-agent scoped fee, or roadmap-and-retainer model? How is ongoing maintenance priced?
5. **vCAIO** — standalone advisory retainer, or layered on top of vCIO?
6. **AI-Ready Managed Services baseline** — included at no uplift (brand moat), or modest uplift vs. non-AI managed services?
7. **Vendor license pass-through** — pass-through at cost, at a fixed markup, or let the client buy direct?
8. **Custom agent license fuel** (Anthropic / OpenAI API usage) — pass-through, bundled, or metered?

> **RECOMMENDATION**
>
> Set these deliberately — in one sitting, with Drew's partners and an advisor — rather than on the fly in front of a prospect. A consistent commercial model is the thing that makes Katana feel like a real business, not a side hustle.

16 Technical Build Specs & the 30-Day Action Plan

Katana’s AI delivery stack

Capability	Tooling	Why
Primary general-purpose model	Anthropic Claude (Sonnet / Opus tiers).	Strongest reasoning for long documents, writing, and agentic workflows.
Secondary model	OpenAI (GPT-series) and Google Gemini available.	Hedge; client-mandated; specific strengths (Gemini for Workspace-native work).
Client productivity AI	Microsoft 365 Copilot or Gemini for Workspace (platform-matched).	Native, secure, defaults for 99% of clients.
Agent framework	Claude Agent SDK for productized agents; direct API for repeat workflows; MCP where justified.	Context- and cost-disciplined.
RAG / knowledge layer	Vendor-native (Copilot / Gemini over Workspace) first; Azure AI Search or Pinecone when custom is justified.	Native beats custom for most SMB cases.
Governance platform	Katana-maintained AUP templates; M365/ Google DLP; audit dashboards.	Reusable across clients.
Internal ops & PSA AI	Existing PSA vendor’s native AI + Katana internal agents.	Runs Katana’s own house.

First 30 days — Katana's AI enablement practice launch

Week	Deliverable	Action
Week 1	Internal readiness.	Lock the offering table (Section 5). Finalize commercial model (Section 15) with Drew and partners. Build the AUP template and readiness-assessment questionnaire.
Week 2	First client-facing asset.	One-pager: "Katana AI Enablement." Deck-light, plain-language, outcomes-led. Based on this primer, not a replacement for it.
Week 3	Pilot readiness assessment.	Run a full 4D assessment against a friendly existing client or partner. Deliver the report. Capture what worked and what didn't.
Week 4	Launch externally.	Update Katana website with AI offering; announce to existing book; line up 5 prospect conversations.

NORTH STAR

By day 90, every Katana managed-services client has been offered a Readiness Assessment, an AUP, and a Copilot or Gemini enablement path. By day 180, the first custom-agent reference case is shippable.

Closing — What This Primer Is and Isn't

This is **version 3** of the AI-for-Clients Primer. It is deliberately opinionated on *what* Katana sells and *how* it delivers — but deliberately silent on Katana's own fees and commercial model, which are Drew's to set.

Next doc in the series: the **AI Enablement One-Pager** — the distilled, client-ready, 1-page pitch derived from this primer. After that, the **4D Readiness Questionnaire** and the **Katana AUP template**.

PREPARED BY

Fructify LLC — *Cultivate growth, engineer success.*
For Drew Fitchett, founder of Katana. April 2026.