FRUCTIFY

Cultivate growth, engineer success.

The Katana AI Readiness Primer

QUESTIONS · OPTIONS · DECISIONS

Prepared for: Drew Fitchett, Founder — Katana

By: Fructify LLC

Date: April 2026

Cultivate growth, engineer success.

© 2026 Fructify LLC · Prepared in partnership with Drew Fitchett / Katana

Table of Contents

01 Where Katana Stands Today

02 The Questions Your Clients Are Already Asking

03 Who’s Really Asking (Role-by-Role Decoder)

04 Where AI Is Actually Landing First (Verticals)

05 “Can You Protect My AI?” — The Governance Entry Point

06 “What Can AI Do For Me?” — The Vendor Landscape

07 The AI Stack — The Building Blocks Behind What Clients Are Asking For

08 Anthropic’s AI Fluency 4D Framework

09 “Do You Offer AI?” — Three Shapes Katana’s Practice Could Take

10 What Setting Up AI for a Client Actually Looks Like

11 Open Questions for Drew to Decide

01 Where Katana Stands Today

Katana today offers the managed IT services Drew has spent nearly three decades building expertise around — the same shape as Anchor Network Solutions: managed endpoints, cybersecurity, cloud, BCDR, vCIO, helpdesk. **Katana does not have an AI practice yet.** That is not a weakness of this document — it is the premise of it.

What has changed is demand. Drew's prospects and existing clients are walking into conversations asking three variations of the same question:

1. *"What can AI actually do for my business?"*
2. *"Do you offer AI solutions?"*
3. *"Can you protect the AI tools we're already using?"*

Each of these is a different conversation with a different answer. This primer is built for Drew — not for his clients — to equip him to navigate all three.

> **WHAT THIS DOCUMENT IS**
>
> A working reference that (a) maps the questions Drew's clients will ask and what's really behind them, (b) surveys the AI landscape those clients are being pitched, (c) lays out the options Drew has for standing up an AI practice, and (d) names the decisions only Drew can make.

> **WHAT THIS DOCUMENT IS NOT**
>
> It is not a service catalog for Katana — Katana has not chosen or built those services yet. It is not a pricing sheet — Katana has not set pricing. It is not a pitch deck for Drew's clients — it is Drew's own reading material.

How to use it

- Read Sections 2–3 before your next AI conversation with a client. The crib sheet and role decoder are immediately usable.
- Read Section 5 ("Can you protect my AI?") closely — it is the most credible thing Drew can answer today because it leans on security, which Katana already delivers.
- Treat Section 7 (three shapes of an AI practice) as a menu of options Drew picks from, not a recommendation.
- Section 9 is a list of decisions only Drew can make. Every one of them is open.

02 The Questions Your Clients Are Already Asking

These are the questions clients are actually walking in with — organized by what they really mean underneath, with a one-line answer Drew can use as a starting point in conversation. Drew can adjust tone and detail in the moment; the point here is to have a confident first sentence ready for each.

Client Question	What They Actually Mean	A Good First Answer from Drew
“Should we be using ChatGPT?”	Are we behind our competitors? Is this a real productivity tool or hype?	Probably yes for some roles, but on a business-tier license (not the free app) so your data stays protected. Let’s look at who on your team would use it and whether Copilot, Gemini, or ChatGPT Business fits your stack best.
“Is Copilot worth it?”	Microsoft is pushing this hard and it’s expensive. Should I buy?	It depends on your M365 tier, how clean your SharePoint/ email data is, and how your team actually works. I’d rather help you pilot it with a small group before you commit a hundred seats.
“Can AI read our email / contracts / tickets?”	I want AI that knows my business without me retyping everything.	Yes — that’s called retrieval. Copilot does it over M365, Gemini does it over Workspace, and custom solutions can do it over anything else. The hard part isn’t the AI, it’s your permissions and data hygiene. We’d look at that first.
“Will AI replace my people?”	I’m anxious. Should I be hiring differently or laying off?	Not in the near term for the work your people actually do. What we see is AI making existing staff faster on the repetitive parts of the job. We can talk about where it would help you most.
“What’s our AI policy?”	My board, insurer, or auditor asked and I don’t have one.	You need one. Most of our clients don’t have one either. We can put a draft in front of you based on NIST’s AI Risk Management Framework and your industry, then roll it out with training.
“Can you make our CRM / ERP smart?”	My business software is clunky. I heard AI can help.	Often yes — through your software vendor’s own AI feature (Salesforce Einstein, HubSpot Breeze, NetSuite AI) or a custom integration. We’d start with what your vendor already offers before building anything custom.
“How do I stop my team from using ChatGPT with company data?”	I know they’re pasting stuff into AI tools. That scares me.	Banning it usually fails. The better move is giving them a sanctioned, licensed tool with data protections, publishing an acceptable-use policy, and adding browser or DLP controls. That’s something we can do for you.
“Is my data safe if we use this?”	What are these AI companies doing with our prompts?	On the business and enterprise tiers of Copilot, ChatGPT, Claude, and Gemini, your prompts and outputs are not used to train the public model. The free consumer versions are different. I’ll document this for each tool you use so you have it in writing.
“Can you protect us while we use AI?”	I’m going to use this anyway. Keep me out of trouble.	Yes — that’s basically the security layer around AI. We set sanctioned tools, an AUP, logging, DLP, and training.

Client Question	What They Actually Mean	A Good First Answer from Drew
		It’s an extension of what I already do for your environment.
“What does this cost?”	I need a number before my next board meeting.	The license side is well-understood — I’ll walk you through vendor pricing. My own fees depend on what we decide to scope. Let’s figure out what you actually need first and I’ll put a number on it.

03 Who’s Really Asking (Role-by-Role Decoder)

The same question means different things from different seats. Reading the room changes what a good answer sounds like.

Role	What's Actually on Their Mind	Where the Conversation Wants to Go
CEO / Owner	Am I behind? Is this a threat or an opportunity?	Talk outcomes and competitive posture. Reference a peer or vertical example. Don't start with tools.
COO / Ops Lead	My people are drowning in busywork.	Talk concrete workflow pain: triage, scheduling, documentation, summaries. Lead with Copilot/Gemini if they already have the license.
CFO / Controller	Another SaaS line item. What's the ROI?	Talk license utilization, hours saved per user, and comparison against the commodity line of the budget. Offer to help rationalize existing AI spend before adding more.
IT Director / IT Manager	AI is landing in every tool — I can't evaluate all of it.	Position as a partner who will do the evaluation work. Offer a stack review. This is where Katana takes load off someone who's drowning.
Security / Compliance Officer	Shadow AI, data leakage, audit readiness.	Lead with governance: sanctioned tools, AUP, logging, DLP, training. This is the conversation where Katana sounds most credible today.
Head of Sales / Marketing	Content velocity, CRM quality, prospecting speed.	Native CRM AI first (Einstein, Breeze, Copilot for Sales). Then talk research and drafting tools.
HR Director	Recruiting volume, policy writing, onboarding content.	Copilot/Gemini for drafting; vendor AI inside the HRIS for recruiting automation.
Line-of-business lead	One specific painful process they own.	Ask what the process is. Map it to an existing vendor capability or flag it as a candidate for a future custom build.

RULE OF THUMB

If the first question is about **tools**, you're probably talking to IT. If it's about **outcomes**, you're talking to the business. If it's about **risk**, you're talking to security or the CEO. Match the register.

04 Where AI Is Actually Landing First (Verticals)

Adoption is not even across industries. A handful of verticals are moving faster than others — which tells Drew where his early AI conversations are most likely to turn into real work.

Vertical	Where AI Is Actually Being Used	Regulatory Notes	Why It Matters for Katana
Professional services (law, accounting, consulting)	Contract and document review, research, client Q&A over a document library, time capture.	Client confidentiality; bar / AICPA rules on AI use.	High willingness to pay; document-heavy workflows map cleanly to Copilot / Gemini.
Healthcare & dental	Clinical documentation, scheduling, patient messaging, coding assistance.	HIPAA — a BAA is required on any AI tool touching PHI.	Drew already handles HIPAA posture. Extending that to AI is a natural fit.
Construction & trades	Estimating, takeoff, RFI drafting, safety documentation, schedule communication.	Lightly regulated; IP on bid content matters.	Procore / STACK / Buildertrend all have AI now. Katana can be the layer that rolls it out.
Financial services (RIA, wealth, small CPA firms)	Client meeting prep, portfolio narrative drafting, compliance review, email triage.	FINRA / SEC / state rules; supervision of communications.	Requires logging and supervision — a security-forward MSP is well-positioned.
Manufacturing & distribution	SOP generation, RFP responses, ERP conversational access, predictive maintenance pilots.	IP protection; shop-floor data sometimes unstructured.	Good fit if client runs NetSuite / Dynamics / Sage; OT-adjacent work rewards caution.
Nonprofits & associations	Grant writing, donor communications, board summaries, event planning.	Donor privacy; communication rules.	Budget-sensitive but loyal; often already on Google Workspace (Gemini bundled).

These are the verticals showing the fastest adoption in MSP territory as of April 2026. It is not an exhaustive list; Drew may have accounts in other verticals where specific use cases are pulling demand forward.

05 “Can You Protect My AI?” — The Governance Entry Point

Of the three questions clients are asking, this is the one Drew can answer most credibly today — because it is an extension of security work Katana already does. Most clients are **already using AI**, often without sanctioned tools and without a policy. They know it. They want help.

What “protecting AI” actually means

Control	What It Does	Where Katana Already Has the Muscle
Sanctioned-tool list	A documented list of the AI tools employees are allowed to use, with data-handling notes for each.	Similar to the approved-software posture Katana already helps clients maintain.
Acceptable Use Policy (AUP)	Plain-English policy: what can and can’t be pasted into AI tools, attribution rules, disclosure, review expectations.	Katana already drafts and updates IT and security policies; AI AUP is a close cousin.
Data classification & DLP	Sensitivity labels on regulated content; DLP rules that block pasting of PII / PHI / PCI into unsanctioned AI tools.	Native to M365 and Google Workspace; Katana configures these today for other use cases.
Identity & access	SSO + MFA on every sanctioned AI tool; deprovisioning when people leave.	Part of every Katana-managed environment.
Logging & audit	Admin-level visibility into who is using which AI tool for what. Required in regulated industries (FINRA supervision, HIPAA audit).	Direct extension of Katana’s SIEM / audit posture.
Training & awareness	Short, role-appropriate training on what to do and not do with AI. Often an insurance or auditor ask.	Same delivery model as Katana’s existing security awareness training.
Vendor review	Before a new AI tool gets sanctioned: data-handling review, BAA / DPA, model-training opt-out, country of processing.	Mirrors Katana’s existing vendor intake for SaaS.
Incident response	Runbook for AI-specific incidents (prompt injection, sensitive-data leakage via AI, harmful output).	Extension of the incident response plan Katana already maintains.

> **WHY THIS MATTERS AS A STARTING POINT**
>
> Drew can honestly say today: “Katana doesn’t sell a custom AI agent yet — but we can absolutely protect and govern the AI you’re already using, and that’s the most urgent thing for most of you right now.” That sentence is true, defensible, and lands a client conversation.

Reference frameworks to cite

- **NIST AI Risk Management Framework (AI RMF 1.0)** — the de facto vendor-neutral governance standard. Useful to name in policy documents and regulated-client conversations.

- **ISO/IEC 42001** — emerging AI management system standard; relevant if a client asks about formal certification.
- **Colorado SB 205** — local AI law (high-risk AI system deployer obligations) — worth being aware of for Colorado-based clients.
- **Sector rules** — HIPAA, FINRA, SEC, PCI DSS each have specific implications for AI use in their regulated domains.

06 "What Can AI Do For Me?" — The Vendor Landscape

When a client asks what AI can do for them, they are usually asking what tools exist and which ones they should consider. This section maps the landscape Drew's clients are being pitched — so he can speak fluently about it without committing Katana to delivering any of it.

> **IMPORTANT — THESE ARE VENDOR PRICES, NOT KATANA'S PRICES**
>
> Every dollar figure below is what the **vendor** (Microsoft, Anthropic, OpenAI, Google, etc.) charges the client directly. **Katana's own fees and commercial model are entirely separate and have not been set.** Vendor pricing changes often — confirm the current number before quoting it in front of a client.

General-purpose AI for business

Product	Who It's For	Vendor List Price (April 2026)	What It's Good At
Microsoft 365 Copilot	M365 E3 / E5 customers.	$30 / user / month, annual commit, on top of M365 license.	Working over email, docs, Teams, SharePoint. Default choice for M365-heavy shops.
Copilot for Sales / Service	Sales and service orgs on Dynamics / Salesforce / HubSpot.	$40–$50 / user / month, on top of M365 Copilot.	CRM-specific workflows, call prep, opportunity notes.
Gemini for Google Workspace	Google Workspace customers.	Bundled into Workspace Business / Enterprise tiers (no separate add-on as of Jan 2025).	Default for Google shops. Strong for document work.
ChatGPT Team / Enterprise	Any org wanting a general-purpose chat AI with enterprise data controls.	Team: $30 / user / month annual. Enterprise: custom negotiated.	Broad-purpose chat and drafting; works well as a second AI alongside Copilot or Gemini.
Anthropic Claude Pro / Team / Enterprise	Orgs valuing reasoning and long-document handling.	Pro: $20 / user / month. Team: $25 / user / month (5 seat min). Enterprise: custom.	Long-document reasoning, writing-heavy work, research.
Perplexity Enterprise Pro	Research-heavy roles (BD, consulting, strategy).	$40 / user / month.	Live-web research with citations.

Native AI inside the software clients already own

Platform	What It Does	How It's Priced
Salesforce Einstein / Agentforce	AI inside Salesforce — summaries, email drafting, agents.	Usage-based (per action / per conversation) plus add-on SKUs; quoted by Salesforce / partner.
HubSpot Breeze	AI inside HubSpot for marketing, sales, and service.	Included in Pro / Enterprise tiers with usage credits.
NetSuite AI / Dynamics Copilot	AI inside the ERP — summaries, narrative reports, conversational GL access.	Varies by module and tier; usually sold by the platform partner.
ConnectWise / Autotask AI	AI inside the PSA — ticket summaries, first-response drafting.	Included in modern tiers or modest add-on.
Vertical SaaS AI	Clio (legal), Procore (construction), Redtail (wealth), Epic (health), etc.	Vendor-set; often bundled; check each platform.

Vertical-specific AI products clients may already be hearing about

Vertical	Representative products
Legal	Harvey, Spellbook, CoCounsel (Thomson Reuters), Clio Duo.
Healthcare	Abridge, Nuance DAX, Suki, Hippocratic AI, Epic's embedded AI.
Finance	Hebbia, AlphaSense, Kensho.
Sales / marketing	Gong AI, Clay, 11x, Regie.ai, Apollo AI, Jasper.
Customer support	Intercom Fin, Zendesk AI, Forethought, Ada.
HR / recruiting	Gem, Paradox, Eightfold, Lattice AI.
Engineering / IT	GitHub Copilot, Cursor, Anthropic's Claude Code, Replit Agent.

Vendor pricing and product availability move quickly. The numbers above are accurate as of April 2026; confirm before quoting.

07 The AI Stack — The Building Blocks Behind What Clients Are Asking For

Drew does not have to *build* every layer of the stack to sell into this market — but he does have to be literate in it. When a client asks “can you do AI for us?” or “how would you connect Claude to our ticketing system?”, the honest answer depends on which layer of this stack the question lives in. This section is a map, not a methodology — it exists so Drew can translate a client request into which pieces need to move.

Layer 1 — The foundation model

The model is the engine. Everything else is plumbing around it. The major families as of April 2026:

Family	Flagship / Tiers	Where it tends to win
Anthropic Claude	Opus (top-tier reasoning) · Sonnet (balanced) · Haiku (fast, cheap)	Long-context document work, agentic workflows, safety-forward clients. Strongest “brand” story around safety and alignment.
OpenAI GPT	GPT flagship · mid · mini tiers, plus “reasoning” variants	Broadest ecosystem, deepest Microsoft integration, widest third-party tool coverage.
Google Gemini	Ultra / Pro / Flash tiers	Workspace-native clients, multimodal (image/video/audio) workloads, long context.
Open-weight / self-hosted	Llama family, Mistral, Qwen, DeepSeek, etc.	Clients with strict data-residency, air-gapped, or cost-sensitive workloads who will run inference themselves.

For most of Katana’s SMB / mid-market clients, this choice is usually made *for* them by their existing ecosystem (M365 → Copilot on GPT; Workspace → Gemini). The model conversation matters most when a client wants a second, non-Microsoft/Google tool, or a custom build.

Layer 2 — How the model is accessed

There are four fundamentally different ways to plug a model into a client’s world. Clients usually don’t know the difference, but the right answer is almost always driven by which of these four shapes fits.

Approach	What it is	When it's right	Build cost for Drew
Platform-native	The AI is already baked into a tool the client owns — Copilot in M365, Gemini in Workspace, Zoom AI Companion, HubSpot Breeze, Salesforce Einstein, ConnectWise Sidekick.	Client just wants "AI in the tools I already use." Fastest ROI.	**Low.** Licensing, configuration, and enablement — no code.
MCP (Model Context Protocol)	Anthropic's open standard for connecting a model to external tools, data, and actions. Broad and growing server ecosystem (Google Drive, GitHub, Linear, Slack, custom in-house).	The client wants Claude (or another MCP-aware agent) to reach into their systems without a custom integration.	**Low–medium.** Often configuration of existing MCP servers; custom servers when a proprietary system is involved.
Direct API	Code calling the model vendor's API directly — Anthropic, OpenAI, Google, AWS Bedrock, Azure OpenAI.	A custom workflow or agent that needs precise control over prompts, context, tools, cost, and latency.	**Medium–high.** Engineering time, hosting, observability, ongoing maintenance.
CLI / script	Anthropic's Claude Code, OpenAI Codex CLI, or a light script that drives a model from a shell or cron job.	Internal automation and one-off "power tools" — report generation, bulk document work, automation that doesn't need a UI.	**Low.** Fast to stand up; best for Drew's own operation or a pilot.

Rule of thumb: start every client conversation at *platform-native* and only move down the stack if there's a specific reason the native tool can't do it. Jumping to a direct API build when Copilot would have worked is how MSPs burn margin.

Layer 3 — Agent frameworks (when a build is actually needed)

If a client has a workflow that doesn't fit an off-the-shelf tool, the build is usually done on top of an agent framework rather than raw API calls. The main ones as of April 2026:

- **Anthropic Claude Agent SDK** — build agents that use tools, hold long context, and run on Claude.
- **OpenAI Assistants / Agents API** — OpenAI's native agent tooling.
- **Microsoft Copilot Studio** — low-code agent builder for M365-heavy shops.
- **Google Agent Builder (Vertex AI)** — Google's equivalent for Workspace-heavy shops.
- **Salesforce Agentforce** — for clients whose system of record is Salesforce.
- **Open-source orchestration** — LangChain, LangGraph, LlamaIndex, Semantic Kernel — used when Drew's engineer (or build partner) needs full control.

Layer 4 — Knowledge and retrieval (RAG)

Most useful client agents need to answer questions over *the client's own content* — SOPs, contracts, policies, ticket history. That requires a retrieval layer:

- **Vendor-native retrieval** — Copilot over SharePoint/OneDrive, Gemini over Drive, ChatGPT "Projects" / Claude "Projects." Usually the right answer for smaller clients.
- **Managed search layers** — Glean, Elastic, Azure AI Search, Amazon Kendra. For mid-market clients with content in multiple systems.
- **Vector databases** — Pinecone, Weaviate, pgvector, Chroma, Azure AI Search. Underlying plumbing when a custom build is warranted.

Layer 5 — Agent patterns Drew is most likely to be asked about

These are the recurring shapes of "can AI do this for us?" requests Katana will hear in 2026. Each is a pattern, not a Katana product — and each one has a native, MCP, or custom version depending on how deep the client wants to go.

Pattern	What it does	Typical first-fit approach
Internal helpdesk bot	Answers employee "how do I…" questions from the client's SOPs / policies.	Copilot Studio or Claude+MCP over SharePoint.
Intake triage	Routes incoming tickets, emails, or forms to the right queue with a drafted first response.	PSA-native (ConnectWise Sidekick) or a Copilot Studio agent.
Knowledge Q&A	"Ask our documents anything" across a content library.	Platform-native (Copilot / Gemini / Claude Projects) before anything custom.
Meeting prep & summary	Pre-reads for upcoming meetings, notes and action items after.	Zoom AI Companion, Copilot in Teams, Gemini in Meet, Granola.
Proposal / document drafter	Produces a first draft of a proposal, SOW, or report from structured inputs.	Copilot for M365, Gemini in Docs, or a custom API build for high-volume shops.
Compliance checker	Reviews a document against a policy or rubric and flags issues.	Custom build; needs tight prompt design and an audit trail.
Sales / account research	Research on a prospect before a call; drafted outreach.	Platform-native (Copilot, Gong, HubSpot Breeze) or Apollo's AI layer.

Layer 6 — Governance and security wrapper

This is the layer Katana is *already* best positioned to deliver on, regardless of which model or approach the client picks. Identity / SSO, data-loss prevention, acceptable-use policy, sensitivity labeling, vendor data-handling review, audit logging. This layer is the same whether the client is using Copilot, Gemini, or a custom agent — and it is where an MSP's existing security muscle translates directly into an AI practice.

Where Katana is likely starting from — and why the rest still matters

Anchor Network Solutions — the reference model Katana is mirroring — is a Microsoft-first MSP. Anchor is a Microsoft Partner with the Small Business Specialist and Small & Midmarket Cloud Solutions Competency designations, runs Microsoft Azure consulting and hosting as a dedicated service line, and delivers Microsoft 365 / Exchange as core offerings. There is no AWS or Google Cloud partnership anywhere on Anchor's site. Because Katana is mirroring that service model, **the most probable default stack for Katana's first clients is the Microsoft path:**

Layer	Katana's likely default (Microsoft)	Equally valid alternatives
Model	GPT family via Azure OpenAI Service; Claude on Azure (via Foundry / Models-as-a-Service).	Anthropic API, OpenAI API, Google Gemini, AWS Bedrock, open-weight self-hosted.
Access approach	Copilot for M365 (platform-native) for most clients.	MCP, direct API, CLI — chosen when a custom build is required.
Agent framework	Microsoft Copilot Studio.	Claude Agent SDK, OpenAI Assistants API, Salesforce Agentforce, Google Agent Builder (Vertex AI), LangChain / LangGraph, Semantic Kernel.
Retrieval	Copilot over SharePoint/OneDrive; Azure AI Search for larger builds.	Glean, Elastic, Amazon Kendra, Pinecone, Weaviate, pgvector, Chroma.
Governance wrapper	Entra ID, Microsoft Purview, Microsoft Defender for Cloud Apps.	Okta / Ping identity, Varonis, Nightfall, Netskope, Zscaler.

WHY THE NON-MICROSOFT PATHS STILL MATTER

Not every client Katana talks to will be a pure Microsoft shop. Some will be Google Workspace all the way down. Some will have AWS-hosted applications even if their productivity suite is Microsoft. Some will already be using Claude Team, ChatGPT Business, or a Salesforce-centric Agentforce build. Drew has to walk into every discovery call able to recognize and work with the stack he actually finds, not the stack Anchor prefers. The Microsoft default is where Katana *starts* — it is not a constraint on what Katana can serve.

TAKEAWAY

Options A and B (from the next section) mostly live at Layers 1, 2-platform-native, and 6. Option C is the only shape that requires Drew to work fluently in Layers 2-MCP / 2-API, 3, and 4. That distinction — not the marketing language — is the real dividing line between the three shapes of practice.

08 Anthropic’s AI Fluency 4D Framework

Anthropic publishes a simple framework for what “using AI well” looks like at the human level. It is called **AI Fluency: The 4Ds — Delegation, Description, Discernment, Diligence.** Drew does not own this framework and does not need to license it — it is published as a course and reference material by Anthropic and can be cited directly. It is included here because it is the cleanest shared vocabulary Drew can hand to a client who is asking “how do I get my people to actually use this well?”

Dimension	The question it answers	What “good” looks like	Typical gap Drew will see
Delegation	When *should* I use AI — and for what?	Users know which tasks are AI-appropriate, which are “AI-with-review,” and which are human-only. Decisions match the stakes.	All-or-nothing — either “AI does everything” or “we don’t trust it.”
Description	How do I tell the AI what I actually want?	Prompts include role, goal, context, constraints, and examples. People think in repeatable prompt patterns, not one-shot questions.	Two-word prompts; expecting the AI to “read their mind.”
Discernment	Is the output any good — and how do I know?	Users check sources, catch hallucinations, notice when a confident answer is actually wrong, verify numbers and citations.	Copy-paste trust; treating AI output as if it has been proofread.
Diligence	Am I using AI *responsibly*?	Users respect data classification, avoid pasting PII / PHI into unsanctioned tools, disclose AI use where required, keep humans in the loop on consequential decisions.	Shadow AI — staff using consumer tools with client data because the sanctioned tool is clunky.

How Drew can use this framework

- **Shared vocabulary.** Gives Drew and a client a clean way to talk about “where are we weak?” without jargon or vendor bias.
- **Readiness scorecard.** On an intake or QBR call, rate the client 1–5 on each D. Three low scores = training and governance are the blockers, not tools.
- **Training frame.** Most generic “AI training” covers Description only. Katana can deliver sessions on Delegation, Discernment, and Diligence — the three Ds most clients never teach.
- **Governance anchor.** Ties naturally into Katana’s existing security posture: Diligence is the DLP / AUP conversation; Discernment is the audit-trail conversation.

> **WHY THIS IS WORTH CITING**
>
> The 4Ds aren’t a Katana framework. They are Anthropic’s — documented publicly, easy for a client to verify, and neutral on vendor choice. That makes them safer to reference in client materials than any proprietary model, and lets Drew point at a credible source when a skeptical client asks “where did this come from?”

09 “Do You Offer AI?” — Three Shapes Katana’s AI Practice Could Take

There is no single “right” AI practice for an MSP. There are three recognizable shapes. Each one answers the client’s “do you offer AI” question differently, carries different build cost for Drew, and protects the business in different ways. None of them is chosen yet.

Option A — Guide & reseller

Answer Drew gives the client: “I help you pick the right AI tool and roll it out. I don’t build custom AI.”

- Help clients pick between Copilot, Gemini, ChatGPT Business, Claude Team.
- Handle licensing, tenant configuration, pilot cohorts, training, and adoption reporting.
- Resell vendor licenses where a partner program exists (Microsoft CSP, Google partner, etc.).

What this requires Drew to build: familiarity with each major vendor’s AI product, a rollout playbook, light training materials. Most of the muscle is already in an MSP’s existing Copilot / Workspace enablement work.

What it does not require: engineers, API experience, custom code, governance frameworks.

Option B — Guide, reseller, and governance

Answer Drew gives the client: “I help you pick the right AI tool, roll it out, and keep you safe while you use it.”

- Everything in Option A, plus:
- AI acceptable-use policy drafting.
- Data-classification and DLP rule configuration around AI tools.
- Ongoing advisory on new AI tools clients are being pitched.
- AI-specific security awareness training.

What this requires Drew to build: a reusable AUP template, a vendor-review checklist, training content. All of this is adjacent to Katana’s existing security work.

What it does not require: engineers, custom code, or taking on custom builds.

Option C — Full AI practice, including custom builds

Answer Drew gives the client: “I help you pick tools, roll them out, keep you safe, and if you have a specific workflow that needs a custom AI agent, I build that too.”

- Everything in Option B, plus:
- Custom AI agents — e.g. intake triage, Q&A over a client’s document library, estimating assistants, compliance checkers.
- Integration between AI tools and client systems (CRM, PSA, ticketing, SharePoint).
- Ongoing maintenance of deployed agents.

What this requires Drew to build: engineering capability (internal hire, fractional partner, or outsourced build partner), API / MCP familiarity, maintenance posture, product thinking. It is a materially different business from Options A and B.

What it does not require: being first. Drew can spend six to twelve months in Option B and enter Option C when the first custom-build client shows up with real budget.

> **HOW TO THINK ABOUT THIS**
>
> Option A is low-risk and fast to stand up, but does not differentiate Katana from any other MSP doing Copilot rollouts. Option B is where most security-forward MSPs are landing in 2026 — defensible, honest, and profitable. Option C is the highest ceiling and also the highest cost / risk to build. Drew does not have to pick today. He does have to know which direction he is pointing.

10 What Setting Up AI for a Client Actually Looks Like

When the client says yes, here is the work that actually happens. This is a generic engagement shape — not a Katana product — and it applies whether Drew ends up in Option A, B, or C.

Phase 1 — Understand the client’s starting point

- What AI tools are they already using (sanctioned or not)?
- What licenses do they already own? (M365 tier, Google Workspace tier, other SaaS.)
- Where does their sensitive data live, and who has access to it?
- Which one or two workflows are they most hoping AI will help?
- Are there regulatory or insurance drivers (HIPAA, FINRA, cyber policy renewal)?

Phase 2 — Get the foundation ready

Foundation Item	What Good Looks Like	What's Usually Broken
Permissions	Least-privilege access on SharePoint / Drive; groups, not individuals; regularly audited.	"Everyone" shares on sensitive folders; former employees still in groups.
Document hygiene	Canonical locations; duplicates cleaned; archive policy.	Multiple versions of the same policy; AI pulls from the wrong one.
Identity	SSO + MFA on every AI tool; deprovisioning pipeline.	Personal accounts used for AI tools; no MFA on everything.
Sensitivity labels / DLP	Labels on PII / PHI / PCI content; DLP rules prevent paste into unsanctioned tools.	No classification at all; nothing stops AI tools from surfacing regulated data.

Phase 3 — Pick and provision the tool

- Match the tool to the client&rsquo>s ecosystem. M365 shop → Copilot. Workspace shop → Gemini. Either can sit alongside ChatGPT Business or Claude Team if there's a specific reason.
- Start with a small pilot cohort — usually 5–20 users — before a full rollout.
- Configure tenant-level settings: data boundaries, retention, admin audit, user settings.
- Verify the business-tier data-handling terms. Document them for the client's file.

Phase 4 — Governance wrapper

- Draft or update the client's AI Acceptable Use Policy.
- Put the approved-tool list in the client's policy library.
- Configure DLP rules that match the client's regulatory regime.
- Deliver short, role-appropriate training — 20–30 minutes is usually enough to start.

Phase 5 — Measure and iterate

- Adoption reporting: who is actually using it, for what, how often.
- Surface a handful of concrete wins to the business sponsor after 30 and 90 days.
- Use QBRs to decide what comes next — expand cohort, add a second tool, or start a custom-build conversation.

11 Open Questions for Drew to Decide

None of these are decisions this document can make for Drew. They are the decisions that turn this primer into an actual business direction.

Positioning

1. Which of the three shapes (A: guide/reseller · B: guide + governance · C: full practice including custom builds) is Drew aiming at for the next 12 months?
2. What is Drew's honest sentence in a client meeting today when someone asks "do you offer AI?" — and when does that sentence change?
3. Does Katana want to lead with "protect your AI" (security-first framing) or "roll out AI" (productivity-first framing) in early marketing?

Partnerships and tooling

1. Which vendor partner programs does Katana join first? Microsoft CSP (for Copilot reselling) is the obvious candidate; Google, Anthropic, and OpenAI each have their own.
2. Does Katana pick a "house" second AI tool alongside Microsoft (most MSPs pick one of Anthropic Claude, OpenAI ChatGPT, or Google Gemini as the go-to when Copilot isn't the right fit)?
3. Which PSA / ticketing AI features does Katana turn on for its own operation (ConnectWise, Autotask)? This is how Katana "eats its own cooking" and can speak from experience.

Skills and staffing

1. Which certifications / trainings does Drew personally pursue first? (Microsoft Copilot specialist, NIST AI RMF familiarity, vendor-specific tracks.)
2. Does Katana need an engineer on staff (or a fractional / outsourced partner) to credibly offer Option C — and if so, when?
3. How is AI folded into Katana's existing helpdesk and vCIO cadence (if at all)?

Commercial

1. How does Katana charge for readiness / assessment work — fixed fee, included in onboarding, or priced per engagement?
2. How does Katana handle AI-tool license pass-through — at cost, with a markup, or let the client buy direct?
3. Is the AI governance wrapper bundled into standard managed services, or a separate line item?
4. For custom builds (if Option C): hourly, fixed project, or retainer?

Go-to-market

1. Which of Katana's existing or pipeline clients gets the first "hey, I want to talk to you about AI" conversation — and when?
2. What is the minimum credible thing Katana puts on the website about AI before the end of Q2 2026?
3. Does Katana want a reference client or case study before going wide with marketing, or lead with a point of view?

> **RECOMMENDED NEXT STEP**
>
> Work through the Positioning block (questions 1–3) first. Everything else follows from those answers. Those three can be answered in a 90-minute session with anyone Drew trusts as a sounding board.

Closing

This primer is the briefing Drew needs before his next three client conversations. It does not commit Katana to anything. It does not pretend Katana has services it hasn't built. It answers the three questions clients are actually asking — and it names the decisions only Drew can make to turn the answers into a business.

> **PREPARED BY**
>
> **Fructify LLC** — *Cultivate growth, engineer success.*
> For Drew Fitchett, founder of Katana. April 2026.